

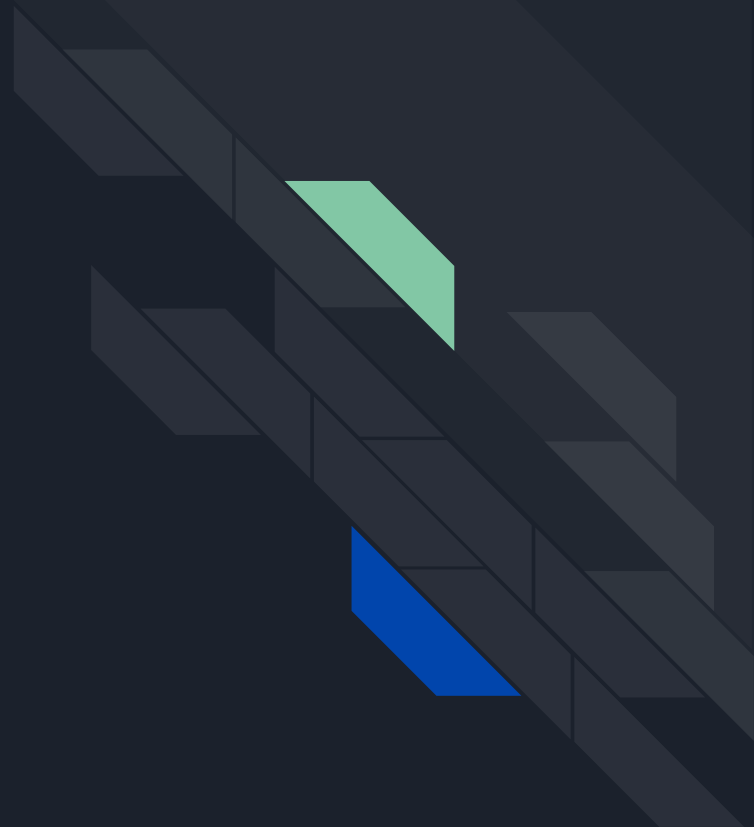
Ingress firewall implementation FreeRTOS

Politecnico di Torino:
Computer Architecture and
Operating Systems course:
2023/2024

Group 5:
Michele , Giorgio , Luca , Gianfranco

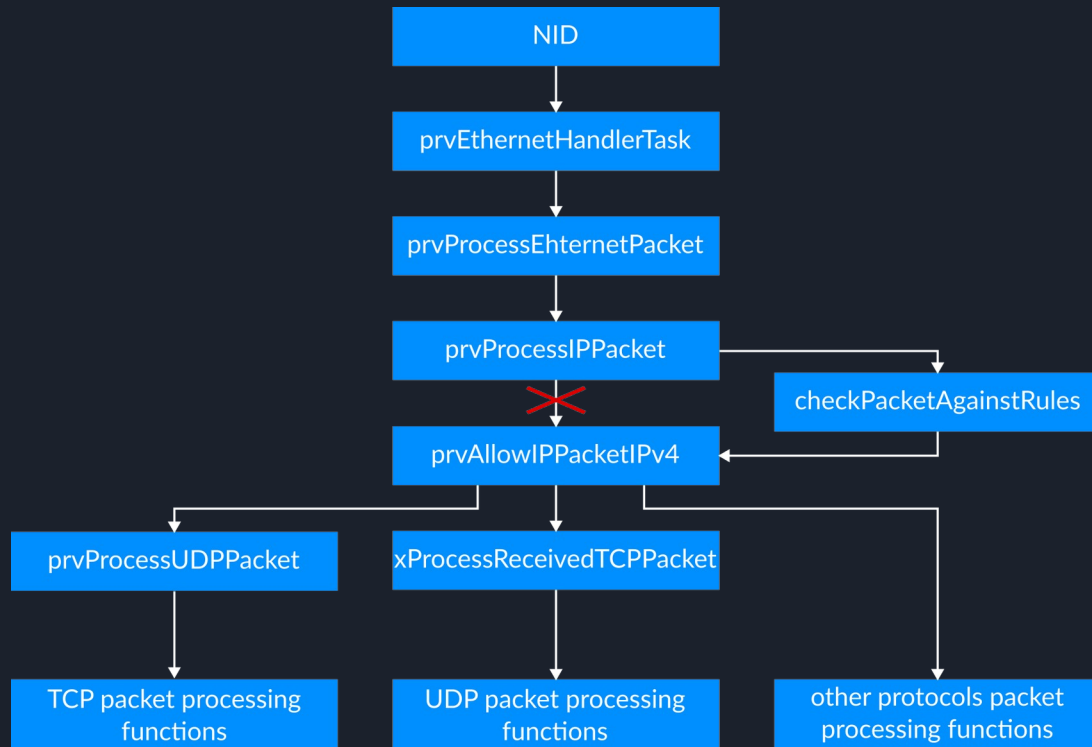
Project intent: create a working firewall for FreeRTOS

- Requirements :
 - Function entry points for filtering packets
 - Handling of rejected packets
 - Print to stderr
 - Rule generation
 - YAML to C struct
 - Rule storage
 - Inline C structs in *rules.h*
 - (WIP) integration with IDS
 - Snort



Understanding FreeRTOS IP stack :

Packet processing pipeline



Functions

```
if( ( uxHeaderLength > ( pxNetworkBuffer ->xDataLength - ipSIZE_OF_ETH_HEADER ) ) ||  
    ( uxHeaderLength < ipSIZE_OF_IPv4_HEADER ) ) || checkPacketAgainstRules (ruleset, NOR, pxIPHeader , pxNetworkBuffer ) )  
  
    // firewall injection point  
    {  
        eReturn = eReleaseBuffer ;  
    }  
}
```

FreeRTOS_IP.c:prvProcessIPPacket

packet interception point

```
uint8_t checkPacketAgainstRules (struct rule ruleset[], int num_rules, const IPHeader_t * pxIPHeader, NetworkBufferDescriptor_t const * pxNetworkBuffer) {  
    //collapse all checks in one function ( fast enough?)  
    switch (pxIPHeader->ucProtocol){  
        case ipPROTOCOL_UDP:{  
            const UDPpacket_t * puUDPpacket = ( ( UDPpacket_t * ) pxNetworkBuffer ->puEthernetBuffer );  
            const UDPHeader_t * puUDPHeader = &( puUDPpacket ->uUDPHeader );  
  
            //Get the header of the UDP packet from the pxNetworkBuffer to extract the ports  
            if (checkIPsAgainstRules (ruleset, num_rules, pxIPHeader, puUDPHeader ->usSourcePort, puUDPHeader ->usDestinationPort) == 1){  
                writeToPcap (pxIPHeader, puUDPHeader ->usSourcePort, puUDPHeader ->usDestinationPort );  
                //If the packet is rejected by the rules pass to login function  
                return 1;  
            }  
            return 0;  
        }  
        break;  
        case ipPROTOCOL_TCP:{  
            const ProtocolHeaders_t *pxProtocolHeaders = ( ( ProtocolHeaders_t * ) &( pxNetworkBuffer ->puEthernetBuffer [ ipSIZE_OF_ETH_HEADER +  
            uxIPHeaderSizePacket ( pxNetworkBuffer ) ] ) );  
            //construct to get to the TCP header fields  
            if (checkPacketsWithPorts (ruleset, num_rules, pxIPHeader, pxProtocolHeaders ->xtCPHeader.usSourcePort, pxProtocolHeaders ->xtCPHeader.usDestinationPort) == 1){  
                writeToPcap (pxIPHeader, pxProtocolHeaders ->xtCPHeader.usSourcePort, pxProtocolHeaders ->xtCPHeader.usDestinationPort );  
                return 1; //return 1 discards the packet  
            }  
            return 0;  
        }  
        break;  
        case ipPROTOCOL_ICMP:{  
            ICMPpacket_t * pxICMPpacket = ( ( ICMPpacket_t * ) pxNetworkBuffer ->puEthernetBuffer ); //ICMP doesn't need ports  
            if (checkIPs (ruleset, num_rules, IPHeader_t * &( pxICMPpacket ->xiPHeader )) == 1){  
                writeToPcap (pxIPHeader, 0, 0);  
                return 1;  
            }  
            return 0;  
        }  
        default:{  
            if (checkIPs (ruleset, num_rules, pxIPHeader) == 1){ //other protocols like ARP  
                writeToPcap (pxIPHeader, 0, 0);  
                return 1;  
            }  
            return 0;  
        }  
    }  
}
```

FreeRTOS_IP.c:checkPacketAgainstRules

```
uint8_t checkIPs (struct rule ruleset[], int num_rules, const IPHeader_t * pxIPHeader){  
    int i;  
    for (i = 0; i < num_rules; i++){  
        if ( (pxIPHeader ->ulSourceIPAddress == ruleset[i].src &&  
            pxIPHeader ->ulDestinationIPAddress == ruleset[i].dst &&  
            pxIPHeader ->ucProtocol == ruleset[i].proto ) {  
            return ruleset[i].action; // match found execute corresponding action on the packet  
        }  
    }  
    return 1; // no match found being this a block list we just let the packet trough  
};  
  
uint8_t checkPacketsWithPorts (struct rule ruleset[], int num_rules, const IPHeader_t * pxIPHeader, uint16_t usSourcePort, uint16_t usDestinationPort) {  
    //collapse all checks in one function ( fast enough?)  
    int i;  
    for (i = 0; i < num_rules; i++){  
        if (pxIPHeader ->ucProtocol == 6){  
            if ( (pxIPHeader ->ulSourceIPAddress == ruleset[i].src &&  
                pxIPHeader ->ulDestinationIPAddress == ruleset[i].dst &&  
                usSourcePort == ruleset[i].port_src && // convert port to network byte order  
                usDestinationPort == ruleset[i].port_dst && // convert port to network byte order  
                pxIPHeader ->ucProtocol == ruleset[i].proto ) {  
                return ruleset[i].action;  
            }  
            // match found execute corresponding action on the packet this can be used both as blacklist and whitelist.  
        }  
        else if (pxIPHeader ->ucProtocol == 17){  
            printf ("%d ", usSourcePort);  
            printf ("%d ", usDestinationPort );  
            if ( (pxIPHeader ->ulSourceIPAddress == ruleset[i].src &&  
                pxIPHeader ->ulDestinationIPAddress == ruleset[i].dst &&  
                usSourcePort == ruleset[i].port_src && // convert port to network byte order -> done in the preprocessor for the rules  
                usDestinationPort == ruleset[i].port_dst && // convert port to network byte order -> done in the preprocessor for the rules  
                pxIPHeader ->ucProtocol == ruleset[i].proto ) {  
                return ruleset[i].action;  
            }  
            // match found execute corresponding action on the packet this can be used both as blacklist and/or whitelist.  
        }  
    }  
    return 1; // no match found being this a allow list we just let the packet trough  
};
```

FreeRTOS_IP.c:checkIPs & checkPacketsWithPort

Handling of rejected packets

```
void writeToPcap(const IPHeader_t * pxIPHeader, uint16_t usSourcePort, uint16_t usDestinationPort){
    //Print to serial console with identifier : clamp fields : and output

    //Write the discarded packets data to stderr -> qemu write output to file -> process in PCAP only the lines with #PCAP -> export in Wireshark
    fprintf(stderr, "#PCAP %d %d %d %d %d \n", pxIPHeader->ulSourceIPAddress, pxIPHeader->ulDestinationIPAddress, pxIPHeader->ucProtocol,
    usSourcePort, usDestinationPort);
}
```

FreeRTOS_IP.c:writeToPcap

- write the data of the rejected packets in *stderr*
 - marking every rejected packet with *#PCAP*
- pipe the *qemu-system-arm* output, that is all directed to *stderr*, to a file out.log
- process the file to compile a *pcap*

```
import socket
import struct
import dpkt
import re

def get_initial_ip(ip_int):
    oct_arr = [(ip_int >> 24) & 255, (ip_int >> 16) & 255, (ip_int >> 8) & 255, ip_int & 255]
    oct_arr.reverse()
    ip_str = ".".join(map(str, oct_arr))
    return ip_str

def get_initial_port(port_int):
    port_int = int(port_int)
    port_int = ((port_int & 0xFF) << 8) | ((port_int >> 8) & 0xFF)
    return port_int

with open('./PART_3/out.log', 'r') as file:
    with open('./output.pcap', 'wb') as pcap_file:
        pcap_writer = dpkt.pcap.Writer(pcap_file)
        for line in file:
            if line.startswith('#PCAP'):
                # Extract the relevant information using regular expressions
                match = re.match(r'#PCAP (Vd+) (Vd+) (Vd+) (Vd+) (Vd+)', line)
                if match:
                    source_ip = get_initial_ip(int(match.group(1)))
                    destination_ip = get_initial_ip(int(match.group(2)))
                    source_port = get_initial_port(int(match.group(4)))
                    destination_port = get_initial_port(int(match.group(5)))
                    protocol = int(match.group(3))

                    eth = dpkt.ethernet.Ethernet()

                    if protocol == 6:
                        tcp = dpkt.tcp.TCP()
                        tcp.sport = source_port
                        tcp.dport = destination_port
                        tcp.data = b'TCP packet'
                        ip = dpkt.ip.IP(src=socket.inet_aton(source_ip), dst=socket.inet_aton(
                            destination_ip), p=dpkt.ip.IP_PROTO_TCP)
                        ip.data = tcp
                        eth.data = ip
                        elif protocol == 17:
                            udp = dpkt.udp.UDP()
                            udp.sport = source_port
                            udp.dport = destination_port
                            udp.data = b'UDP packet'
                            ip = dpkt.ip.IP(src=socket.inet_aton(source_ip), dst=socket.inet_aton(
                                destination_ip), p=dpkt.ip.IP_PROTO_UDP)
                            ip.data = udp
                            eth.data = ip
                        elif protocol == 1:
                            # Create an ICMP packet
                            icmp = dpkt.icmp.ICMP()
                            icmp.type = dpkt.icmp.ICMP_ECHO
                            icmp.data = b'ICMP packet'
                            ip = dpkt.ip.IP(src=socket.inet_aton(source_ip), dst=socket.inet_aton(
                                destination_ip), p=dpkt.ip.IP_PROTO_ICMP)
                            ip.data = icmp
                            eth.data = ip
                        else:
                            ip = dpkt.ip.IP(src=socket.inet_aton(source_ip), dst=socket.inet_aton(
                                destination_ip))
                            eth.data = ip

                        pcap_writer.writepkt(eth)

        # Print a message to indicate the completion of the task
        print('Packets written to output.pcap file.')
```

Rules generation

```
ruleset:  
- source: 192.168.122.50  
  destination: 192.168.122.10  
  port_source: 200  
  port_destination: 4050  
  protocol: 17  
  action: 0  
- source: 192.168.122.1  
  destination: 192.168.122.10  
  port_source: 200  
  port_destination: 33  
  protocol: 6  
  action: 0  
- source: 192.168.122.1  
  destination: 192.168.122.10  
  port_source: ANY  
  port_destination: ANY  
  protocol: 1  
  action: 0
```

rules.yaml

Rules storage

```
#ifndef __RULE_SET_FIREWALL_  
#define __RULE_SET_FIREWALL_ 1  
  
#include <stdlib.h>  
  
#define NOR 3  
  
//A Firewall rule structure  
typedef struct rule {  
    uint32_t src; // Source IP address in network byte order  
    uint32_t dst; // Destination IP address in network byte order  
    uint16_t port_src; // Source port number in network byte order  
    uint16_t port_dst; // Destination port number in network byte order  
    uint8_t proto; // 2-bit mask representing protocol type  
    uint8_t action;  
    //action to do with packets from the source ip 0 >> accept packets 1 >> reject packets  
}Rule;  
  
//Current ruleset loaded in the firewall  
Rule ruleset[NOR] = {  
    {846899392, 175810752, 51200, 53775, 17, 0 },  
    {24815808, 175810752, 51200, 8448, 6, 0 },  
    {24815808, 175810752, 0, 0, 1, 0 },  
};  
  
#endif
```

FreeRTOS-Plus-TCP/rules.h:

Tests for rules and output PCAP

```
send(IP(dst="192.168.122.10" , src="192.168.122.50" )/UDP(dport=4050 , sport=200)/Raw(load="abc" ),
iface="virbr0" )
send(IP(dst="192.168.122.10" , src="192.168.122.1" )/TCP(dport=33 , sport=200), iface="virbr0" )
send(IP(dst="192.168.122.10" , src="192.168.122.1" )/ICMP(), iface="virbr0" )
send(IP(dst="192.168.122.10" , src="192.168.122.50" )/UDP(dport=3050 , sport=200)/Raw(load="def" ),
iface="virbr0" )
send(IP(dst="192.168.122.10" , src="192.168.122.1" )/TCP(dport=6968 , sport=200), iface="virbr0" )
send(IP(dst="192.168.122.10" , src="192.168.122.50" )/ICMP(), iface="virbr0" )
```

Packet generation calls that are fed to **scapy** : 6 packets 3 of which are expected to be rejected by the rules defined earlier

Apply a display filter ... <Ctrl-/>					
No.	Time	Source	Destination	Protocol	Length Info
1	0.000000	192.168.122.50	192.168.122.10	UDP	52 200 → 3050 Len=0
2	0.000055	192.168.122.1	192.168.122.10	TCP	64 200 → 6968 [SYN] Seq=0 Win=65535 Len=10
3	0.000140	192.168.122.50	192.168.122.10	ICMP	49 Echo (ping) request id=0x4943, seq=19792/20557, ttl=64 (no response found!)

The 3 rejected packets as expected

Integration of IDS

- preliminary work for POC using SNORT to process rejected packets by the firewall
- needs second NIC or SNORT endpoint with TCP retransmission mechanism

```
scccccp///pSP///p      p//Y |      -- Jean De Clerck
sY/////////y caa      S//P |
cayCyayP//Ya      pY/Ya
sY/PsY///YCc      aC//Yp
sc sccaCY//PCypaapyCP//YSs
      spCPY/////////YPSps
      ccaacs

using IPython 8.5.0
send(IP(dst="192.168.122.10",src="192.168.122.50")/UDP(dport=4050,sport=200)
/Raw(load="abc"), iface="virbr0")
send(IP(dst="192.168.122.10",src="192.168.122.1")/TCP(dport=33,sport=200),
iface="virbr0")
send(IP(dst="192.168.122.10",src="192.168.122.1")/ICMP(), iface="virbr0")
send(IP(dst="192.168.122.10",src="192.168.122.50")/UDP(dport=3050,sport=200)
/Raw(load="def"), iface="virbr0")
send(IP(dst="192.168.122.10",src="192.168.122.1")/TCP(dport=6968,sport=200)
, iface="virbr0")
send(IP(dst="192.168.122.10",src="192.168.122.50")/ICMP(), iface="virbr0")
```

```
gi4n@ubuntu:~$ sudo snort -q -A console -c /etc/snort/rules/local.rules -i virbr0
02/25-13:36:25.430810  [**] [1:10000001:0] ICMP Traffic Detected [**] [Priority: 0]
CMP} 192.168.122.1 -> 192.168.122.10
02/25-13:36:25.431324  [**] [1:10000001:0] ICMP Traffic Detected [**] [Priority: 0]
CMP} 192.168.122.10 -> 192.168.122.1
02/25-13:36:25.529350  [**] [1:10000001:0] ICMP Traffic Detected [**] [Priority: 0]
CMP} 192.168.122.50 -> 192.168.122.10
```

Rules defined to monitoring only the ICMP packets through SNORT, only proof of concept needs more work

Thanks !

And now time for the live demo 🙌

